


Deadlocks

The Deadlock Problem

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set
- Example1
 - ➔ System has 2 tape drives
 - ➔ P_1 and P_2 each hold one tape drive and each needs another one
- Example2
 - ➔ Semaphores A and B , initialized to 1

| | |
|------------|-----------|
| P_0 | P_1 |
| $Wait(A);$ | $Wait(B)$ |
| $Wait(B);$ | $Wait(A)$ |

Bridge Crossing Example



- Traffic only in one direction.
- Each section of a bridge can be viewed as a resource.
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback).
- Several cars may have to be backed up if a deadlock occurs.
- Starvation is possible.

System Model

- Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- Each resource type R_i has W_i instances
- Each process utilizing a resource:
 - Makes Request for resource
 - Uses the resource
 - Releases the resource
- Request and release of resources is done through system calls

Deadlock Characterization

Deadlock can arise if following **four** conditions hold simultaneously.

- **Mutual exclusion:** only one process at a time can use a resource(i.e. resource is non-sharable).
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes.
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that

P_0 is waiting for a resource that is held by P_1

P_1 is waiting for a resource that is held by P_2

\dots , P_{n-1} is waiting for a resource that is held by P_n

and P_n is waiting for a resource that is held by P_0 .


Resource-allocation Graph

A set of vertices V and a set of edges E .

- V is partitioned into two types:
 - ➔ $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system.
 - ➔ $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system
- Request edge – directed edge $P_i \rightarrow R_j$
- Assignment edge – directed edge $R_j \rightarrow P_i$

Resource-Allocation Graph (Cont.)


- Process




- Resource Type with 4 instances



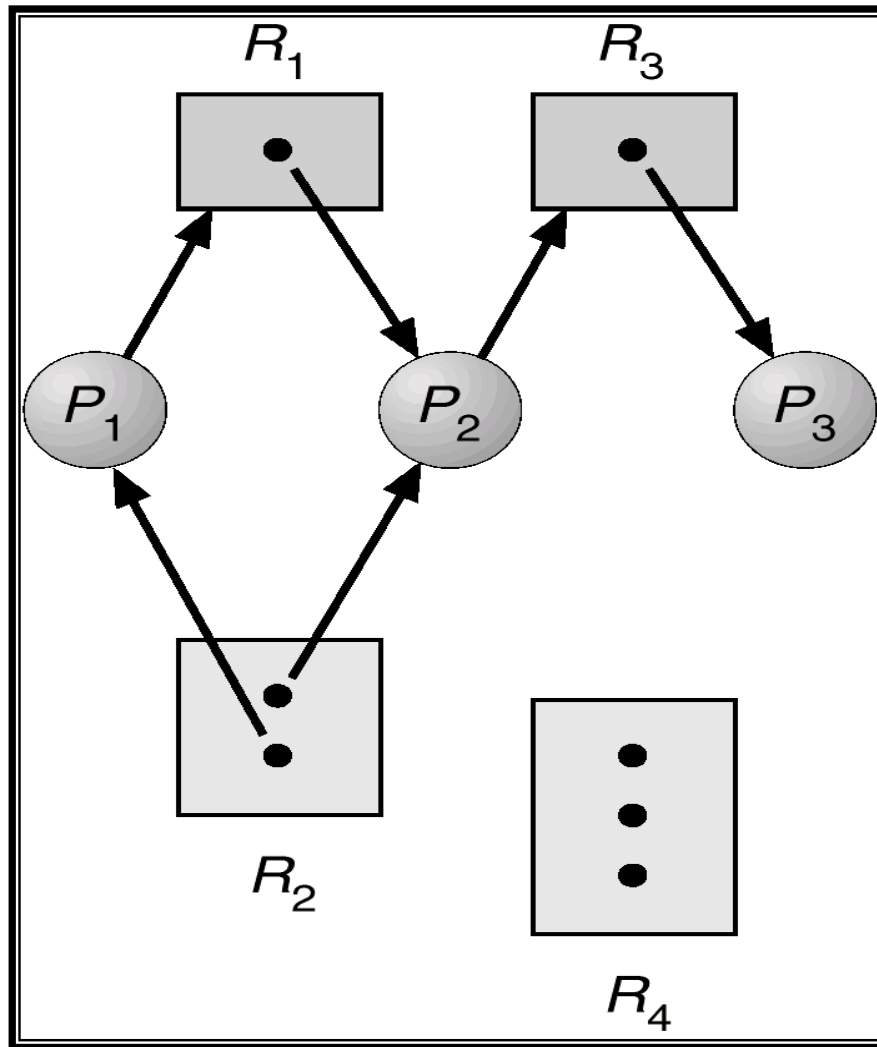
- P_i requests instance of R_j



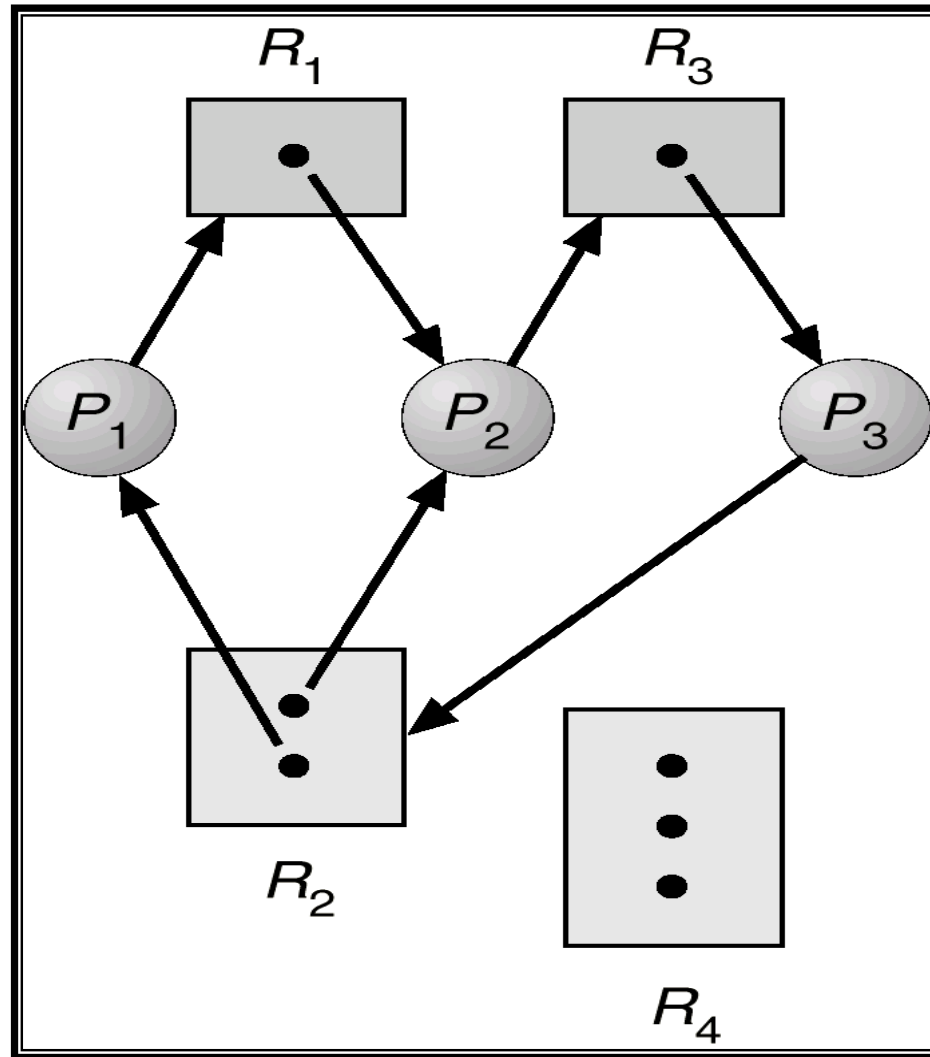
- P_i is holding an instance of R_j



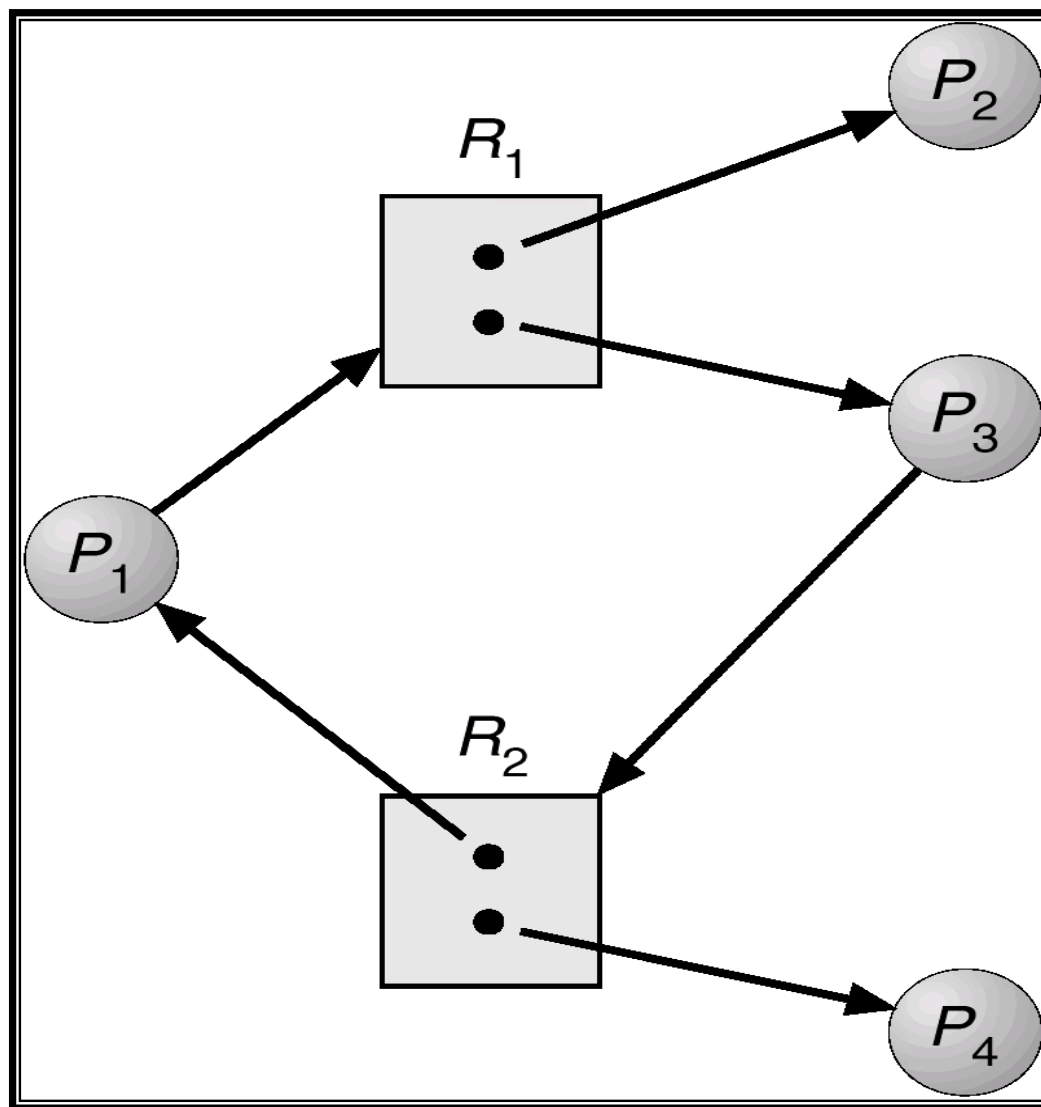
Example of a Resource Allocation Graph



Resource Allocation Graph With A Deadlock



Resource Allocation Graph With A Cycle But No Deadlock



Basic Facts

- If graph contains no cycles \Rightarrow no deadlock.
- If graph contains a cycle \Rightarrow
 - \rightarrow If only one instance per resource type, then deadlock
 - \rightarrow If several instances per resource type, possibility of deadlock

Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state.
(Prevention or Avoidance)
- Allow the system to enter a deadlock state and then Detect & recover.

Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

- **Mutual exclusion**
- **Hold and wait**
- **No preemption**
- **Circular wait**

Deadlock Prevention

Ensure that one of the 4 necessary conditions for deadlock does not hold.

Constraints the way requests can be made for resources

- **Mutual exclusion**

- not required for sharable resources;
- Must hold for non-sharable resources.

➔ Make non-sharable resources as sharable

➔ some resources are intrinsically non sharable

Deadlock Prevention

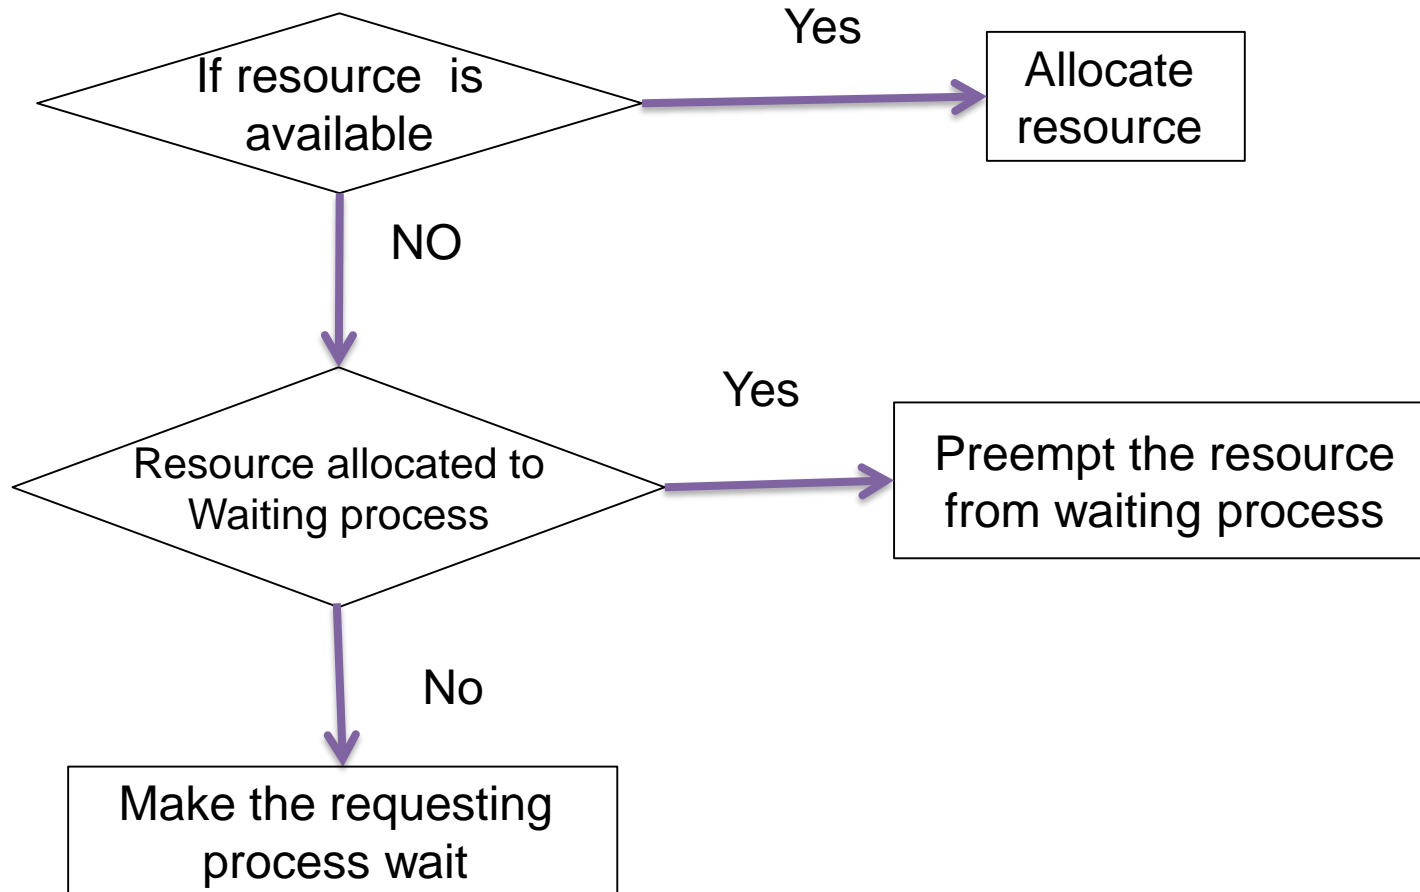
- **Hold and wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
 - ➔ Require process to request and be allocated all its resources before it begins execution,
 - Need to know all the requirement in the beginning itself
 - ➔ Allow process to request resources only when the process has none.
 - If a process has resource and requires additional resources, it must release all resources it is holding and then make request
 - ➔ Low resource utilization; Starvation possible

Deadlock Prevention (Cont.)

- **No preemption – Approach -1**

- ➔ If a process that is holding some resources requests another resource that cannot be immediately allocated to it,
 - Release all resources currently being held by process
- ➔ Preempted resources are added to the list of resources for which the process is waiting
- ➔ Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.
- ➔ Cannot be generally applied to resources like printers, tape drives etc.

No preemption (cont-)



Deadlock Prevention (Cont.)

- **Circular wait** – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration

Deadlock Avoidance

Requires that the system has some additional *a priori* information available.

- Simplest model requires that each process declare the *maximum number* of resources of each type that it may need.
- The deadlock-avoidance algorithm
 - Dynamically finds the *resource-allocation state* in order to *ensure that a circular-wait condition does not occur*.
- Resource-allocation *state* is defined by
 - the number of available resources
 - the number of allocated resources,
 - the maximum requirement of the processes

Safe State

- When a process requests an available resource, the system needs to decide if immediate allocation leaves the system in a *safe state*.
- A state is safe if system can allocate resources to each process (Max requirement) in some order and still avoid deadlock
- System is in safe state if there exists a *safe sequence* of all processes.

- Sequence $\langle P_1, P_2, \dots, P_n \rangle$ is safe if for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$
 - ➔ If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished
 - ➔ When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
 - ➔ When P_i terminates, P_{i+1} can obtain its needed resources, and so on

Example


| | Maximum need | Allocation |
|----|--------------|------------|
| P0 | 10 | 5 |
| P1 | 4 | 2 |
| P2 | 9 | 2 |

Total no of resources =12
safe sequence < P1,P0,P2>

Basic Facts

- If a system is in safe state \Rightarrow no deadlocks.
- If a system is in unsafe state \Rightarrow possibility of deadlock.
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state


Safe, Unsafe , Deadlock State




Resource-Allocation Graph Algorithm

- *Claim edge* $P_i \rightarrow R_j$ indicates that process P_i may request resource R_j ; represented by a dashed line.
- Claim edge converts to *request edge* when a process requests a resource.
- When a resource is released by a process, *assignment edge* reconverts to a claim edge.
- Resources must be claimed *a priori* in the system.

Resource-Allocation Graph For Deadlock Avoidance



Unsafe State In Resource-Allocation Graph



Banker's Algorithm

- Multiple instances.
- Each process must a priori claim maximum use.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.

Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- **Available:** Vector of length m . If $Available[j] = k$, there are k instances of resource type R_j available.
- **Max:** $n \times m$ matrix. If $Max[i, j] = k$, then process P_i may request at most k instances of resource type R_j .
- **Allocation:** $n \times m$ matrix. If $Allocation[i, j] = k$ then P_i is currently allocated k instances of R_j .
- **Need:** $n \times m$ matrix. If $Need[i, j] = k$, then P_i may need k more instances of R_j to complete its task.

$$Need[i, j] = Max[i, j] - Allocation[i, j]$$

Safety Algorithm

1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively.
Initialize:
$$Work = Available$$
$$Finish[i] = false \text{ for } i = 0, 1, \dots, n-1.$$
2. Find an index *i* such that both:
(a) $Finish[i] == false$
(b) $Need_i \leq Work$
If no such *i* exists, go to step 4.
3. $Work = Work + Allocation_i$
 $Finish[i] = true$
go to step 2.
4. If $Finish[i] == true$ for all *i*, then the system is in a safe state.

Safe state example

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 3 | 2 | 2 |
| P2 | 6 | 1 | 3 |
| P3 | 3 | 1 | 4 |
| P4 | 4 | 2 | 2 |

Claim Matrix

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 1 | 0 | 0 |
| P2 | 6 | 1 | 2 |
| P3 | 2 | 1 | 1 |
| P4 | 0 | 0 | 2 |

Allocation Matrix

| R1 | R2 | R3 |
|----|----|----|
| 9 | 3 | 6 |

Resource Vector

| R1 | R2 | R3 |
|----|----|----|
| 0 | 1 | 1 |

Available Vector

(a) Initial state

| | R1 | R2 | R3 |
|----|----|----|----|
| P1 | 2 | 2 | 2 |
| P2 | 0 | 0 | 1 |
| P3 | 1 | 0 | 3 |
| P4 | 4 | 2 | 0 |

Resource-Request Algorithm for Process P_i

$Request_i$ = request vector for process P_i .

If $Request_i[j] = k$ then process P_i wants k instances of resource type R_j .

1. If $Request_i \leq Need_i$ then go to step 2.
Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If $Request_i \leq Available$, go to step 3.
Otherwise P_i must wait, since resources are not available.
3. Pretend to allocate requested resources to P_i by modifying the state as follows:
 $Available = Available - Request_i;$
 $Allocation_i = Allocation_i + Request_i;$
 $Need_i = Need_i - Request_i$
 - If safe \Rightarrow the resources are allocated to P_i .
 - If unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored

Example of Banker's Algorithm

- 5 processes P_0 through P_4 ; 3 resource types A (10 instances), B (5 instances) and C (7 instances).
- Snapshot at time T_0 :

| | <u>Allocation</u> | | | <u>Max</u> | | | <u>Available</u> | | |
|-------|-------------------|---|---|------------|---|---|------------------|---|---|
| | A | B | C | A | B | C | A | B | C |
| P_0 | 0 | 1 | 0 | 7 | 5 | 3 | 3 | 3 | 2 |
| P_1 | 2 | 0 | 0 | 3 | 2 | 2 | | | |
| P_2 | 3 | 0 | 2 | 9 | 0 | 2 | | | |
| P_3 | 2 | 1 | 1 | 2 | 2 | 2 | | | |
| P_4 | 0 | 0 | 2 | 4 | 3 | 3 | | | |

Example (Cont.)

- The content of the matrix Need is equal to Max – Allocation.

| | <u>Need</u> | | |
|-------|-------------|----------|----------|
| | <i>A</i> | <i>B</i> | <i>C</i> |
| P_0 | 7 | 4 | 3 |
| P_1 | 1 | 2 | 2 |
| P_2 | 6 | 0 | 0 |
| P_3 | 0 | 1 | 1 |
| P_4 | 4 | 3 | 1 |

| | <u>Allocation</u> | <u>Max Available</u> | | |
|-------|-------------------|----------------------|----------|-------|
| | <i>A</i> | <i>B</i> | <i>C</i> | |
| P_0 | 0 | 1 | 0 | 7 5 3 |
| P_1 | 2 | 0 | 0 | 3 2 2 |
| P_2 | 3 | 0 | 2 | 9 0 2 |
| P_3 | 2 | 1 | 1 | 2 2 2 |
| P_4 | 0 | 0 | 2 | 4 3 3 |

- The system is in a safe state since the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies safety criteria.

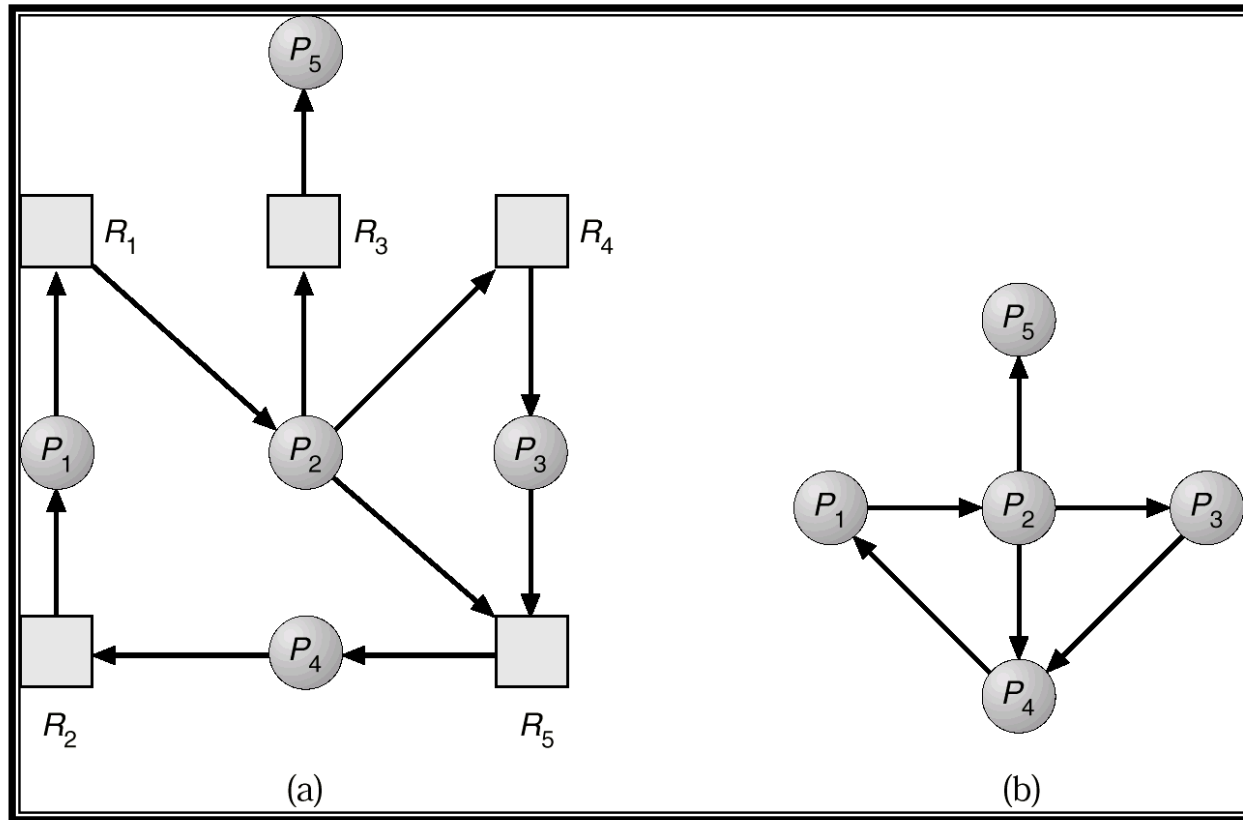
Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme

Single Instance of Each Resource Type

- Maintain *wait-for* graph
 - ➔ Nodes are processes
 - ➔ $P_i \rightarrow P_j$ if P_i is waiting for P_j
- Periodically invoke an algorithm that searches for a cycle in the graph.

Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph

Corresponding wait-for graph

Several Instances of a Resource Type

- *Available*: A vector of length m indicates the number of available resources of each type.
- *Allocation*: An $n \times m$ matrix defines the number of resources of each type currently allocated to each process.
- *Request*: An $n \times m$ matrix indicates the current request of each process. If $Request[i,j] = k$, then process P_i is requesting k more instances of resource type R_j .

Detection Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively. Initialize:
 - (a) *Work* = *Available*
 - (b) For $i = 1, 2, \dots, n$, if $Allocation_i \neq 0$, then $Finish[i] = false$; otherwise, $Finish[i] = true$.
2. Find an index i such that both:
 - (a) $Finish[i] == false$
 - (b) $Request_i \leq Work$If no such i exists, go to step 4.
3. $Work = Work + Allocation_i$
 $Finish[i] = true$
go to step 2.
4. If $Finish[i] == false$, for some i , $1 \leq i \leq n$, then the system is in deadlock state. Moreover, if $Finish[i] == false$, then P_i is deadlocked.

Example of Detection Algorithm

- Five processes P_0 through P_4 ; three resource types
A (7 instances), B (2 instances), and C (6 instances).
- Snapshot at time T_0 :

| | <i>Allocation</i> | | | <i>Request</i> | | | <i>Available</i> | | |
|-------|-------------------|----------|----------|----------------|----------|----------|------------------|----------|----------|
| | <i>A</i> | <i>B</i> | <i>C</i> | <i>A</i> | <i>B</i> | <i>C</i> | <i>A</i> | <i>B</i> | <i>C</i> |
| P_0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| P_1 | 2 | 0 | 0 | 2 | 0 | 2 | | | |
| P_2 | 3 | 0 | 3 | 0 | 0 | 0 | | | |
| P_3 | 2 | 1 | 1 | 1 | 0 | 0 | | | |
| P_4 | 0 | 0 | 2 | 0 | 0 | 2 | | | |

- Sequence $\langle P_0, P_2, P_3, P_4, P_1 \rangle$ will result in $Finish[i] = \text{true}$ for all i .

Example (Cont.)

- P_2 requests an additional instance of type C .

Request

| | A | B | C |
|-------|-----|-----|-----|
| P_0 | 0 | 0 | 0 |
| P_1 | 2 | 0 | 1 |
| P_2 | 0 | 0 | 1 |
| P_3 | 1 | 0 | 0 |
| P_4 | 0 | 0 | 2 |

- State of system?
 - ➔ Can reclaim resources held by process P_0 , but insufficient resources to fulfill other processes requests.
 - ➔ Deadlock exists, consisting of processes P_1 , P_2 , P_3 , and P_4 .

Detection-Algorithm Usage

- When, and how often, to invoke depends on:
 - ➔ How often a deadlock is likely to occur?
 - ➔ How many processes will need to be rolled back?
- Frequency of execution
 - Invoke the algorithm every time request for resource cannot be granted immediately.
 - Once per hour
 - When CPU utilization drops below 40%

Recovery from Deadlock: Process Termination

- Abort all deadlocked processes → May be expensive
- Abort one process at a time until the deadlock cycle is eliminated → Invoke deadlock detection algorithm after every abort.
- In which order should we choose to abort?
 - Priority of the process.
 - How long process has computed, and how much longer to completion.
 - Resources the process has used.
 - Resources process needs to complete.
 - Is process interactive or batch?

Recovery from Deadlock: Resource Preemption

- Selecting a victim – minimize cost.
- Rollback – return to some safe state, restart process for that state.
- Starvation – same process may always be picked as victim, include number of rollback in cost factor.