

# Virtual Memory

# Background

- **Virtual memory** – separation of user logical memory from physical memory.
  - ➔ Only part of the program needs to be in memory for execution.
  - ➔ Logical address space can therefore be much larger than physical address space.
  - ➔ Allows address spaces to be shared by several processes.
  - ➔ Allows for more efficient process creation.
- Virtual memory can be implemented via:
  - ➔ Demand paging
  - ➔ Demand segmentation

# Demand Paging

- Bring a page into memory only when it is needed.
  - ➔ Less I/O needed
  - ➔ Less memory needed
  - ➔ Faster response
  - ➔ More users
- Page is needed  $\Rightarrow$  reference to it
  - ➔ invalid reference  $\Rightarrow$  abort
  - ➔ not-in-memory  $\Rightarrow$  bring to memory

# Valid-Invalid Bit

- With each page table entry a valid–invalid bit is associated  
(1  $\Rightarrow$  in-memory, 0  $\Rightarrow$  not-in-memory)
- Initially valid–invalid bit is set to 0 on all entries.
- Example of a page table snapshot.

Frame #	valid-invalid bit
	1
	1
	1
	1
	0
⋮	
	0
	0


page table

- During address translation, if valid–invalid bit in page table entry is 0  $\Rightarrow$  page fault.

# Page Fault

- If there is ever a reference to a page, first reference will trap to OS  $\Rightarrow$  page fault
- OS looks at another table to decide:
  - $\rightarrow$  Invalid reference  $\Rightarrow$  abort.
  - $\rightarrow$  Just not in memory.
- Get empty frame.
- Swap page into frame.
- Reset tables, validation bit = 1.
- Restart instruction:

# Steps in Handling a Page Fault



# What happens if there is no free frame?

- Page replacement – find some page in memory, but not really in use, swap it out.
  - ➔ algorithm
  - ➔ performance – want an algorithm which will result in minimum number of page faults.
- Same page may be brought into memory several times.

# Performance of Demand Paging

- Page Fault Rate  $0 \leq p \leq 1.0$

- if  $p = 0$  no page faults

- if  $p = 1$ , every reference is a fault

- Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p (\text{page fault overhead}) \end{aligned}$$

[swap page out + swap page in + restart overhead]



# Demand Paging Example

- Memory access time = 1 microsecond
- 50% of the time the page that is being replaced has been modified and therefore needs to be swapped out.
- Swap Page Time = 10 msec = 10,000 microsec

$$\text{EAT} = (1 - p) \times 1 + p (10000)$$
$$1 + 10000p$$

# Page Replacement

- Page-fault service routine includes page replacement.
- Use *modify (dirty) bit* to reduce overhead of page transfers – only modified pages are written to disk.
- Page replacement completes separation between logical memory and physical memory
  - large virtual memory can be provided on a smaller physical memory.

# Replacement Policy

- Which page to replaced?
- Page removed should be the page least likely to be referenced in the near future
- Most policies predict the future behavior on the basis of past behavior

# Replacement Policy


- Frame Locking
  - ➔ If frame is locked, it may not be replaced
  - ➔ Kernel of the operating system
  - ➔ Key control structures
  - ➔ I/O buffers
  - ➔ Associate a lock bit with each frame

# Page Replacement Algorithms

- Want lowest page-fault rate.
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string.
- In all our examples, the reference string is  
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5.

# First-In-First-Out (FIFO) Algorithm


- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (3 pages can be in memory at a time per process)




# First-In-First-Out (FIFO) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 4 frames (4 pages can be in memory at a time )
- In general more frames  $\Rightarrow$  less page faults
- FIFO Replacement – Belady's Anomaly

10 page faults



# FIFO Illustrating Belady's Anomaly






# Optimal Algorithm

- Replace page that will not be used for longest period of time.
- 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5




6 page faults

- Used for measuring how well algorithm performs.

# Least Recently Used (LRU) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# LRU Implementation

## ■ Counter implementation

- ➔ Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter.
- ➔ When a page needs to be changed, look at the counters to determine which are to change.

# LRU Algorithm (Cont.)

- Stack implementation – keep a stack of page numbers in a double link form:
  - Page referenced:
    - move it to the top
  - No search for replacement

No Belady's anomaly → Stack Algorithms

# LRU Approximation Algorithms

## ■ Reference bit

- With each page associate a bit, initially = 0
- When page is referenced bit set to 1.
- Replace the one which is 0 (if one exists). We do not know the order, however.


## ■ Second chance

- Need reference bit.
- If page to be replaced (in clock wise order) has reference bit = 1. then:
  - set reference bit 0.
  - leave page in memory.
  - replace next page (in clock wise order), subject to same rules.

# The Clock Policy

- A method to give 'a chance' to recently used pages
  - ➔ a *new* page is not replaced unless there is no other choice
- The set of frames candidate for replacement is considered as a circular buffer
- When a page is replaced, a pointer is set to point to the next frame in buffer
- A use bit for each frame is set to 1 whenever
  - ➔ a page is first loaded into the frame
  - ➔ the corresponding page is referenced
- When it is time to replace a page, the first frame encountered with the use bit = 0 is replaced.
  - ➔ During the search for replacement, each use bit set to 1 is changed to 0

# Second-Chance (clock) Page-Replacement Algorithm




# Enhanced Clock Policy

- In addition to reference bit use modify bit also
  - ➡ (0,0) not referenced not modified
  - ➡ (0, 1) Not recently used but modified
  - ➡ (1,0) recently used but not modified
  - ➡ (1,1) recently used and modified



# Comparison



**Figure 8.17** Comparison of Fixed-Allocation, Local Page Replacement Algorithms

# Counting Algorithms

- Keep a counter of the number of references that have been made to each page.
- LFU Algorithm: replaces page with smallest count.
- MFU Algorithm: based on the argument that the page with the smallest count was probably just brought in and has yet to be used.

# Allocation of Frames

- Each process needs **minimum** number of pages.
- Example:
- MOV source, destination
  - instruction is 4 bytes, might span 2 pages.
  - 2 pages to handle **from**.
  - 2 pages to handle **to**.
- Two major allocation schemes.
  - fixed allocation
  - priority allocation

# Fixed Allocation

- Equal allocation – e.g., if 100 frames and 5 processes, give each 20 pages.
- Proportional allocation – Allocate according to the size of process.

$s_i$  = size of process  $p_i$

$$S = \sum s_i$$

$m$  = total number of frames

$$a_i = \text{allocation for } p_i = \frac{s_i}{S} \times m$$

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_2 = \frac{127}{137} \times 64 \approx 59$$

# Priority Allocation

- Use a proportional allocation scheme using priorities rather than size.
- If process  $P_i$  generates a page fault,
  - ➔ select for replacement one of its frames.
  - ➔ select for replacement a frame from a process with lower priority number.


# Global vs. Local Allocation

- **Global** replacement – process selects a replacement frame from the set of all frames; one process can take a frame from another.
  - ➔ Process cannot control its own Page fault rate
- **Local** replacement – each process selects from only its own set of allocated frames.
  - ➔ Number of frames allocated to a process do not change
  - ➔ Does not make use of less used pages belonging to other processes

# Thrashing

- If a process does not have “enough” pages, the page-fault rate is very high. This leads to:
  - ➔ low CPU utilization.
  - ➔ operating system thinks that it needs to increase the degree of multiprogramming.
  - ➔ another process added to the system.
- **Thrashing**  $\equiv$  a process is busy swapping pages in and out.
  - ➔ More pronounced for Global page replacement policy

# Thrashing



- Why does paging work?  
Locality model
  - ➔ Process migrates from one locality to another.
  - ➔ Localities may overlap.
- Why does thrashing occur?  
 $\Sigma$  size of locality > total memory size




# Working-Set Model

- $\Delta \equiv$  working-set window  $\equiv$  a fixed number of page references  
Example: 10,000 instruction
- $WSS_i$  (working set of Process  $P_i$ ) = total number of pages referenced in the most recent  $\Delta$  (varies in time)
  - if  $\Delta$  too small will not encompass entire locality.
  - if  $\Delta$  too large will encompass several localities.
  - if  $\Delta = \infty \Rightarrow$  will encompass entire program.
- $D = \sum WSS_i \equiv$  total demand frames
- if  $D > m$  (*Total number of available frames*)  $\Rightarrow$  Thrashing
- Policy if  $D > m$ , then suspend one of the processes.


# Working-set model

page reference table

. . . 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 . . .




$WS(t_1) = \{1, 2, 5, 6, 7\}$



$WS(t_2) = \{3, 4\}$

# Page-Fault Frequency Scheme



- Establish “acceptable” page-fault rate.
  - ➔ If actual rate too low, process loses frame.
  - ➔ If actual rate too high, process gains frame.

# Other Considerations

- Page Buffering:

- ➔ Maintain a pool of free frames to quickly restart a faulting process
- ➔ Can be used to improve performance of some simple page replacement algorithms like FIFO

- Prepaging:

- ➔ Bring in the complete working set of a swapped out process to avoid initial multiple faults

# Other Considerations (Cont.)

- **TLB Reach** - The amount of memory accessible from the TLB.
- $\text{TLB Reach} = (\text{TLB Size}) \times (\text{Page Size})$
- Ideally, the working set of each process is stored in the TLB. Otherwise there is a high degree of page faults.

# Other Considerations (Cont.)

## ■ Program structure

→ `int A[ ][ ] = new int[1024][1024];`

→ Each row is stored in one page

→ Program 1

```
for (j = 0; j < A.length; j++)  
  for (i = 0; i < A.length; i++)  
    A[i,j] = 0;
```

1024 x 1024 page faults

→ Program 2

```
for (i = 0; i < A.length; i++)  
  for (j = 0; j < A.length; j++)  
    A[i,j] = 0;
```

1024 page faults