

innovate

achieve

lead



**BITS Pilani**  
Pilani Campus

# Operating Systems

Computer Science and Information Systems Department  
BITS, Pilani



# Inter Process Communication

## Tutorial 9

# Today's Agenda

---



- Inter Process Communication
  - Pipes
  - FIFOs
  - Message Queues
  - Shared Memory

# Introduction



- Inter process communication (IPC) is a mechanism which allows processes to communicate each other. This involves synchronizing their actions and managing shared data.
- Communication can be of two types –
  - Between related processes initiating from only one process, such as parent and child processes.
  - Between unrelated processes, or two or more different processes.

# Introduction contd.

---

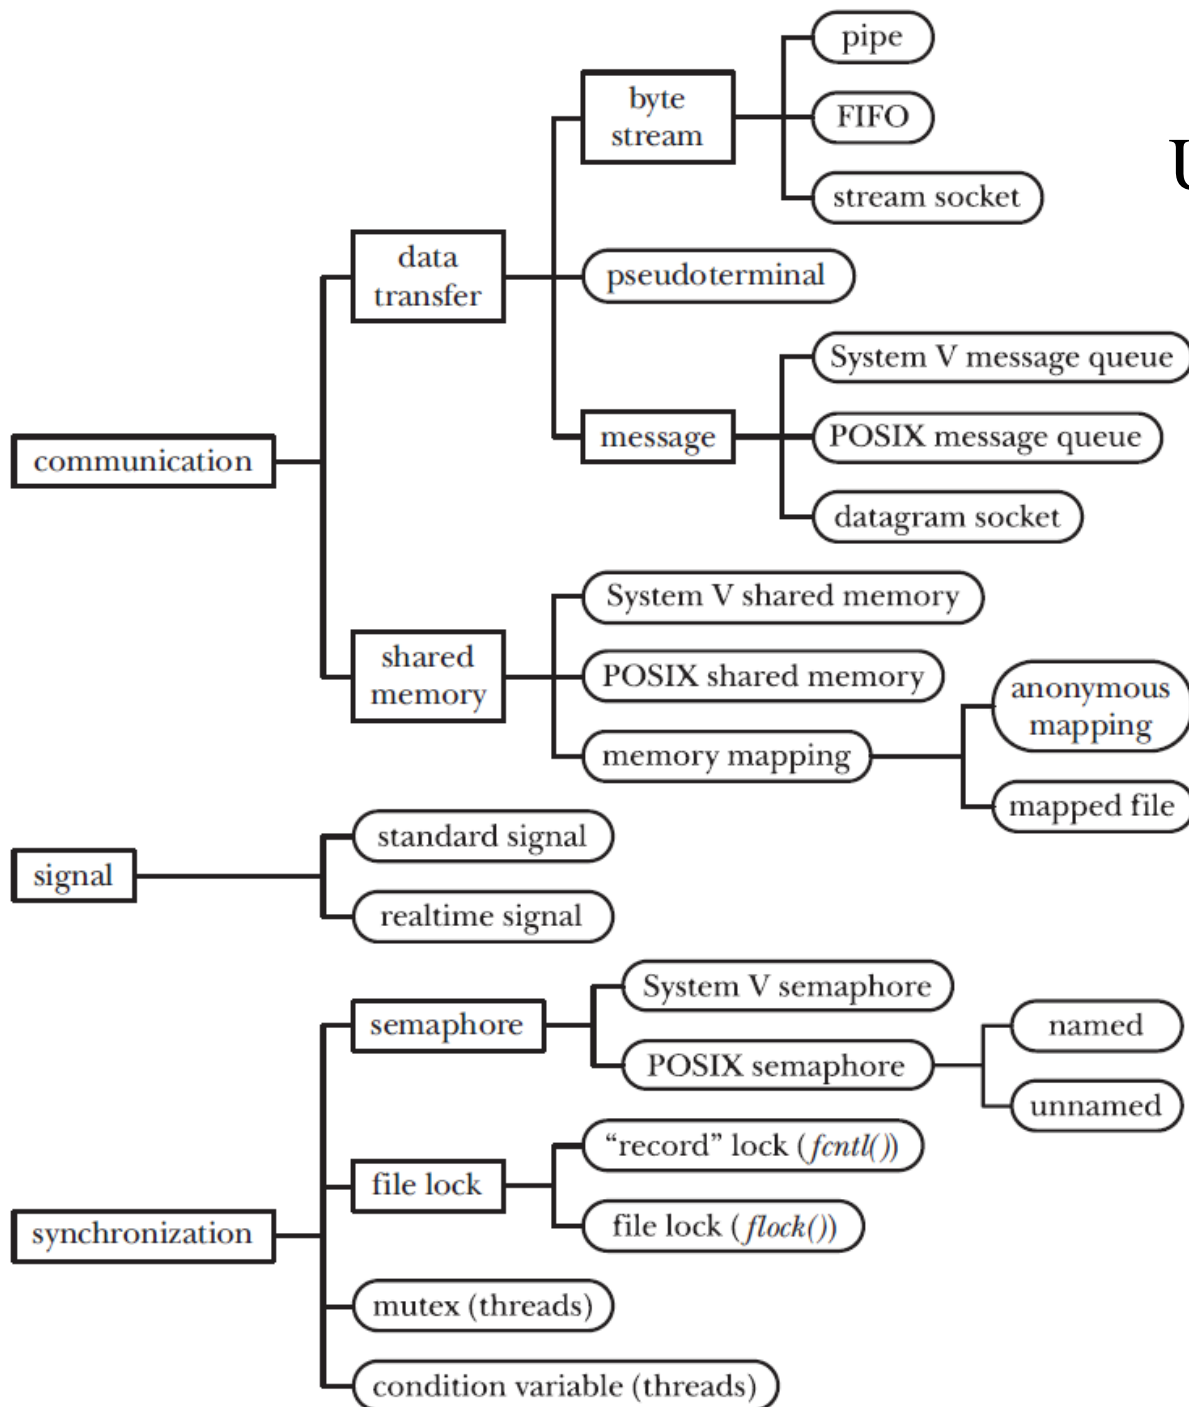
- **Pipes**– Communication between two **related processes**. The mechanism is **half duplex** meaning the first process communicates with the second process. To achieve a full duplex i.e., for the second process to communicate with the first process another pipe is required.
- **FIFO**– Communication between two **unrelated processes**. FIFO is a **full duplex**, meaning the first process can communicate with the second process and vice versa at the same time.
- **Message Queues**– Communication between two or more processes with **full duplex** capacity. The processes will communicate with each other by posting a message and retrieving it out of the **queue**.

# Introduction contd.

---


- **Shared Memory**– Communication between two or more processes is achieved through a **shared piece of memory** among all processes.
- **Semaphores**– Semaphores are meant for **synchronizing** access to multiple processes. When one process wants to access the memory (for reading or writing), it needs to be **locked** (or protected) and **released** when the access is removed.
- **Signals**– Signal is a mechanism to communication between multiple processes by way of **signaling**. This means a source process will send a signal (recognized by number) and the destination process will handle it accordingly.

# A taxonomy of UNIX IPC facilities



- **Data Transfer Facility:**

- In order to communicate, one process **writes** data to the IPC facility, and another process **reads** the data.
- These facilities require two data transfers between user memory and kernel memory: one transfer from user memory to kernel memory during writing, and another transfer from kernel memory to user memory during reading.





# Communization (Data transfer Categories)



## – Byte stream:

- The data exchanged via pipes, FIFOs and datagram sockets is an **undelimited** byte stream.
- Each read operation may read an **arbitrary** number of bytes from the IPC facility, regardless of the size of blocks written by the writer.

## – Message:

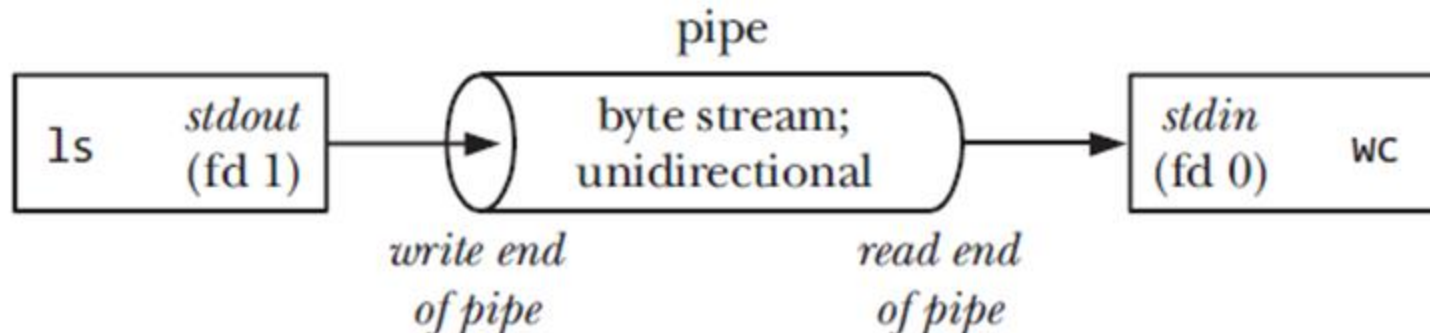
- The data exchanged message queues
- These are form of **delimited** messages.
- Each read operation reads **entire message**,
- It is not possible to read part of a message, leaving the remainder on the IPC facility; nor is it possible to read multiple messages in a single read operation.

- Shared Memory

- Shared memory allows processes to exchange information by placing it in a region of memory that is **shared** between the processes
- A process can make data available to other processes by placing it in the shared memory region. Because communication **doesn't require system calls** or **data transfer** between user memory and kernel memory, shared memory can provide **very fast** communication.
- There is a need to **synchronize operations** on the shared memory.
- For example, one process should not attempt to access a data structure in the shared memory while another process is updating it.
- A **semaphore** is the usual synchronization method used with shared memory.

# Pipes (Byte Stream)

- Shell Command
  - `$ ls | wc -l`
- To execute the above command, the shell creates two processes, executing `ls` and `wc`, respectively using `fork()` and `exec()` system calls.
- Two processes are connected to the pipe so that the **writing process** (`ls`) has its standard output (file descriptor 1) joined to the write end of the pipe, while the **reading process** (`wc`) has its standard input (file descriptor 0) joined to the read end of the pipe.
- Pipes are unidirectional



# Creating and Using Pipes

---



## Syntax:

```
#include <unistd.h>
```

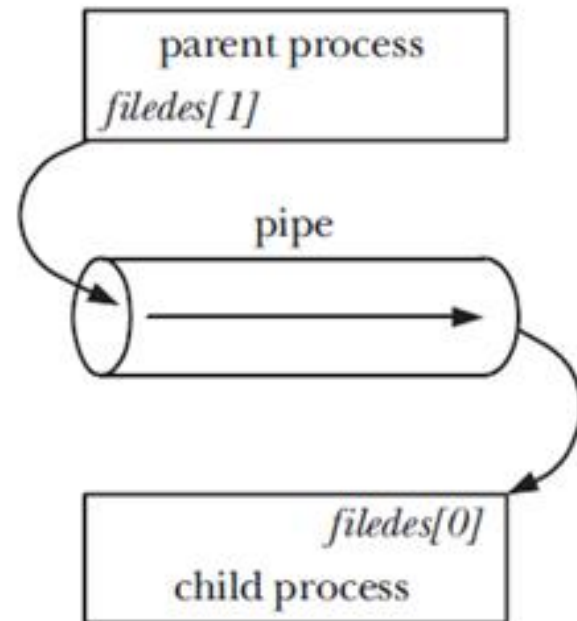
```
int pipe(int filedes[2]);
```

- Returns 0 on success, or  $-1$  on error
- A successful call to `pipe()` returns two open file descriptors in the array `filedes`: one for the read end of the pipe (`filedes[0]`) and one for the write end (`filedes[1]`).
- we can use the `read()` and `write()` system calls to perform I/O on the pipe.

# Creating and Using Pipes



a) After *fork()*



b) After closing unused descriptors

# **PIPE BETWEEN PARENT AND CHILD PROCESS**

```
int main()
{
    int filedес[2];
    if (pipe(filedes) == -1) printf("error creating pipe \n");
    else
        printf("Pipe Created Successfully\nfiledes[0] = %d,filedes[1] = %d\n",filedes[0],filedes[1]);
    switch (fork()) { /* Create a child process */
    case -1:
        printf("fork failed \n");
    case 0: /* Child */
        printf("Child Process.....%d\n",getpid());
        if (close(filedes[1]) == -1)    printf("failed to close \n");
        else
            printf("Closed write end, read file descr = %d\n",filedes[0]);
        break;
    default: /* Parent */
        wait(NULL);
        printf("Parent Process.....%d\n",getpid());
        if (close(filedes[0]) == -1)    printf("failed to close \n");
        else
            printf("Closed read end, write file descr = %d\n",filedes[1]);
        break; } }
```

```

int main()
{
    int filedес[2];
    if (pipe(filedes) == -1) printf("error creating pipe \n");
    else
        printf("Pipe Created Successfully\nfiledes[0] = %d,filedes[1] = %d\n",filedes[0],filedes[1]);
    switch (fork()) { /* Create a child process */
    case -1:
        printf("fork failed \n");
    case 0: /* Child */
        printf("Child Process.....%d\n",getpid());
        if (close(filedes[1]) == -1) printf("failed to close \n");
        else
            printf("Closed write end, read file descr = %d\n",filedes[0]);
        break;
    default: /* Parent */
        wait(NULL);
        printf("Parent Process.....%d\n",getpid());
        if (close(filedes[0]) == -1) printf("failed to close \n");
        else
            printf("Closed read end, write file descr = %d\n",filedes[1]);
        break; } }

```

```

$ ./a.out
Pipe Created Successfully
filedes[0] = 3,filedes[1] = 4
Child Process.....1966
Closed write end, read file
descr = 3
Parent Process.....1965
Closed read end, write file
descr = 4

```



```
int main(void)
{
    int pfd[2];
    char buf[30];

    pipe(pfd);

    if (!fork()) {
        printf(" CHILD: writing to the pipe\n");
        write(pfd[1], "test", 5);
        printf(" CHILD: exiting\n");
        exit(0);
    }
    else {
        printf("PARENT: reading from pipe\n");
        read(pfd[0], buf, 5);
        printf("PARENT: read \"%s\"\n", buf);
        wait(NULL);
    }

    return 0;
}
```

PARENT: reading from pipe  
CHILD: writing to the pipe  
CHILD: exiting  
PARENT: read "test"

```
Int main(int argc, char *argv[])
```

```
{
```

```
    int pfd[2];
```

```
    char buf[BUF_SIZE];
```

```
    ssize_t numRead;
```

```
    if (argc != 2 || strcmp(argv[1], "--help") == 0)
```

```
        printf("%s %s \n", argv[0], argv[1]);
```

```
    if (pipe(pfd) == -1)
```

```
        printf("failed to create pipe\n");
```

```
    else
```

```
        printf("Pipe created successfully\nread file des = %d, write file des = %d\n", pfd[0], pfd[1]);
```

```
    switch (fork()) {
```

```
    case -1:
```

```
        printf("failed to fork\n");
```

```
    case 0:
```

```
        /* Child - reads from pipe */
```

```
        printf("Child Created.....%d\n", getpid());
```

```
        numRead = read(pfd[0], buf, BUF_SIZE);
```

```
        printf("text read from pipe by child : %s , numRead = %d\n", buf, numRead);
```

```
        if (numRead == -1)            printf("read error\n");
```

```
        if (numRead == 0)            break; /* End-of-file */
```

```
        if (write(STDOUT, buf, numRead) != numRead)
```

```
            printf("child - partial/failed write\n");
```

```
        write(STDOUT, "\n", 1);
```

```
        if (close(pfd[0]) == -1) printf("failed to close");
```

```
            exit(EXIT_SUCCESS);
```

```
    default:
```

```
        /* Parent - writes to pipe */
```

```
        printf("Parent continues %d\n", getpid());
```

```
        if (close(pfd[0]) == -1)        /* Read end is unused */
```

```
            printf("failed to close read end in - parent\n");
```

```
        if (write(pfd[1], argv[1], strlen(argv[1])) != strlen(argv[1]))
```

```
            printf("parent - partial/failed write\n");
```

```
        else
```

```
            printf("Parent successfully write to pipe (%d): %s\n", pfd[1], argv[1]);
```

```
        if (close(pfd[1]) == -1) printf("failed to close");
```

```
        wait(NULL);
```

```
        exit(EXIT_SUCCESS);
```

```
    } }
```

**Using a pipe to communicate  
between a parent and child  
process**

**\$ ./a.out BITS**

**Pipe created successfully**

**read file des = 3, write file des = 4**

**Parent successfully write to pipe**

**(4): BITS**

**Child Created.....1750**

**text read from pipe by child : BITS ,**

**numRead = 4**

**BITS**

# Usage of Pipes to execute commands like `ls | wc`



```
int pfd[2];  
pipe(pfd);    /* Allocates (say) file descriptors 3 and 4 for pipe */  
  
/* Other steps here, e.g., fork() */  
  
close(STDOUT);    /* Free file descriptor 1 */  
dup(pfd[1]);       /* Duplication uses lowest free file descriptor, i.e.,  
                  fd 1 */  
  
dup2(pfd[1], STDOUT); /* Close descriptor 1, and reopen bound to write  
                     end of pipe */
```

# Programming Exercises

---



- Write a C program to implement the unix command “ls -l | sort” using pipes.

---

Any Queries?